

# Streaming and Massively Parallel Algorithms for Edge Coloring

Soheil Behnezhad, Mahsa Derakhshan, Marina Knittel,  
MohammadTaghi Hajiaghayi, **Hamed Saleh** [ESA'19]



# Massively Parallel Computation

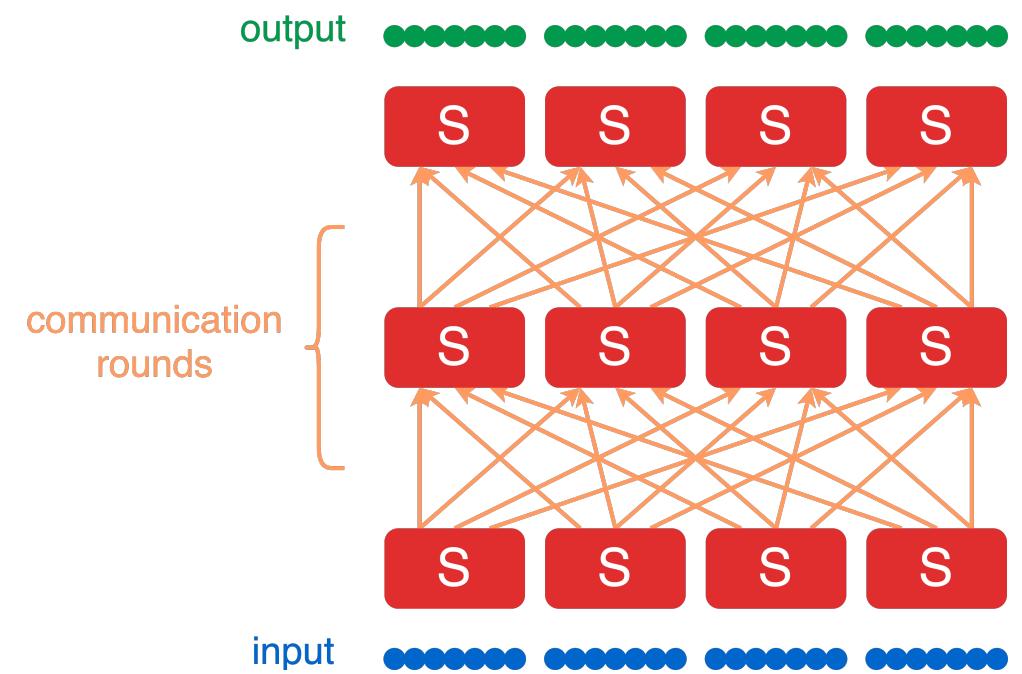
- Modern frameworks for Large-scale parallel/distributed data processing: MapReduce, Hadoop, Spark.
- Key Idea: distribute the workload among several machines.
- The **MPC** model: A theoretical model to abstract out the computational power of these frameworks.

[Karloff et. al 2010]  
[Goodrich et. al 2011]  
[Beame et. al 2013]  
[Andoni et. al 2014]



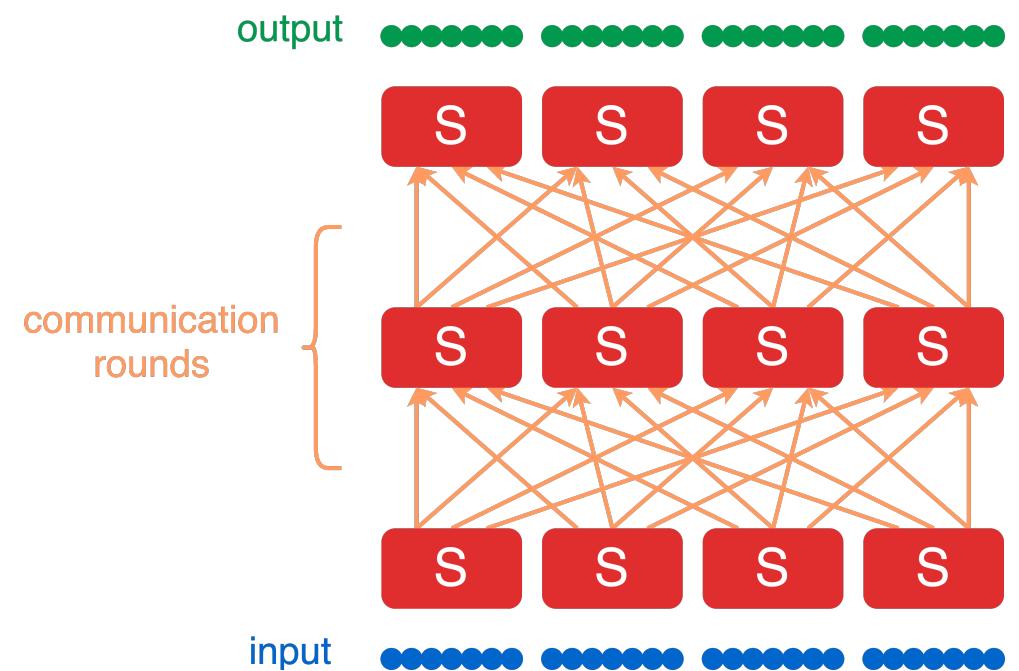
# Massively Parallel Computation

- There are  $M$  machines each with memory  $S$ .
- The input of length  $N$  is initially (randomly) distributed among the machines.
  - Sublinear space  $S = o(N)$ .
  - Usually  $N = O(S \cdot M)$ .
- The data is processed in several **synchronous rounds**.



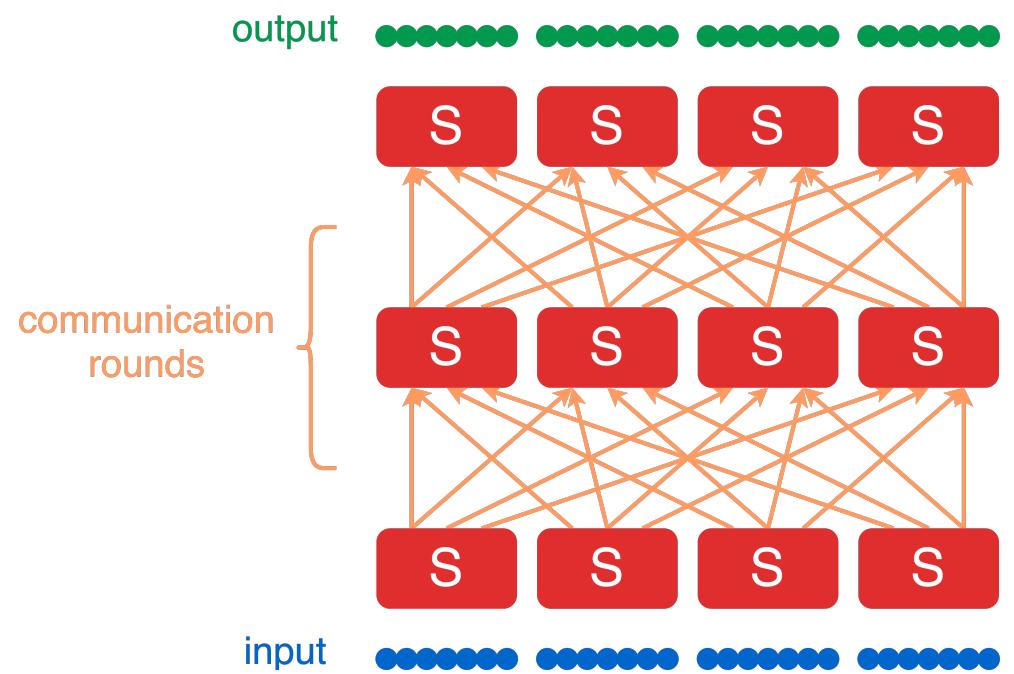
# Massively Parallel Computation

- There are  $M$  machines each with memory  $S$ .
- The data is processed in several **synchronous rounds**. In each round,
  - Machines perform arbitrary computation on their **local** data.
  - Machines **communicate** with each other. Total **incoming/outgoing** messages of each machine is bounded by  $O(S)$  words.



# Massively Parallel Computation

- There are  $M$  machines each with memory  $S$ .
- The data is processed in several **synchronous rounds**.
- Main bottleneck: **Communication**. We wish for algorithms with very small number of rounds.
  - Often **sub-logarithmic** rounds.

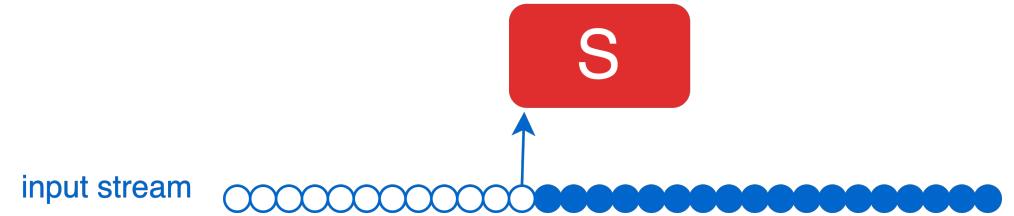


# Related **distributed/parallel** models

- The **PRAM** model (shared memory)
  - Any PRAM algorithm running in time  $t = t(n)$  can be simulated in  $\mathbf{O}(t)$  MPC rounds.  
[Karloff et. al 2010]
- The **LOCAL**, **CONGEST**, and **congested-clique** models
  - Similar techniques can be used for both MPC and these distributed models.
  - Congested-clique is almost equivalent to MPC (in terms of the number of rounds).  
[Behnezhad et. al 2018]

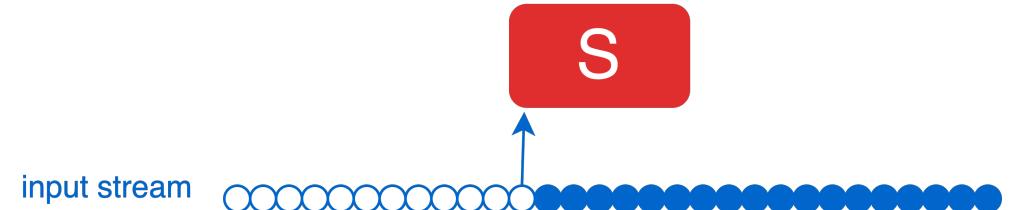
# Streaming

- There is a single machine with memory  $\textcolor{brown}{S}$ .
- The input of length  $\textcolor{brown}{N}$  is streamed into the machine.
  - Sublinear space  $\textcolor{brown}{S} = o(\textcolor{brown}{N})$ .
- The data is processed in several **passes**.



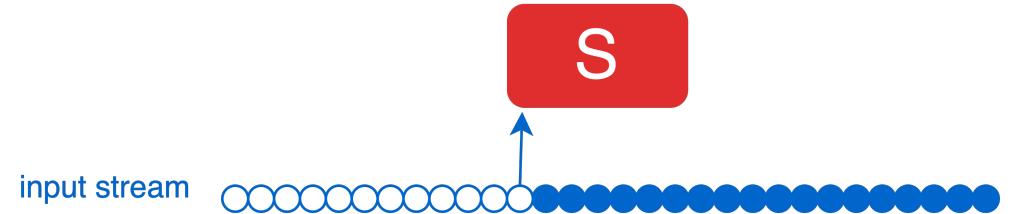
# Streaming

- There is a single machine with memory **S**.
- The data is processed in several **passes**. In each pass:
  - The input entries arrive **sequentially** and **one by one** in a specific order.
  - The order can be either **random** or **adversarial**.



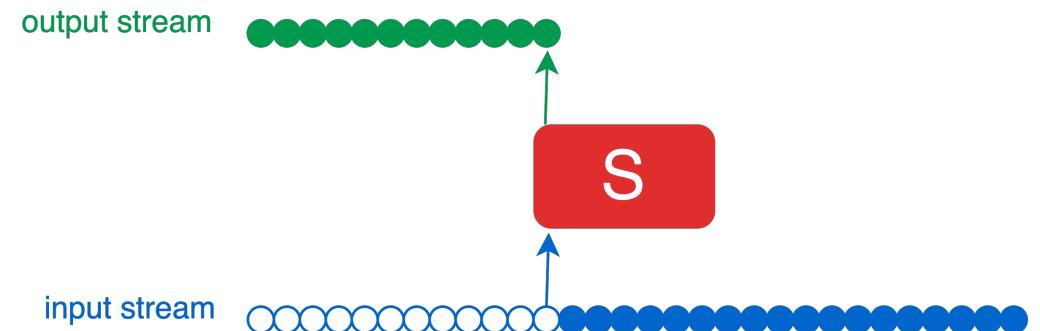
# Streaming

- There is a single machine with memory **S**.
- The data is processed in several **passes**.
- There is a **trade-off** between the number of passes and the space of the machine.



# W-streaming

- There is a single machine with memory  $S$ .
- What if the **output** also doesn't fit in the memory?
  - We use the **W-streaming** model.
  - The output is also streamed.



# Semi-streaming and Semi-MPC

- In many graph problems, we assume  $S = O(|V|)$ , where the input graph is given as  $G = (V, E)$ .
- The variant of the streaming model in which  $S = O(|V|)$  is called **Semi**-streaming.

[Feigenbaum et. al 2005]  
[McGregor 2014]

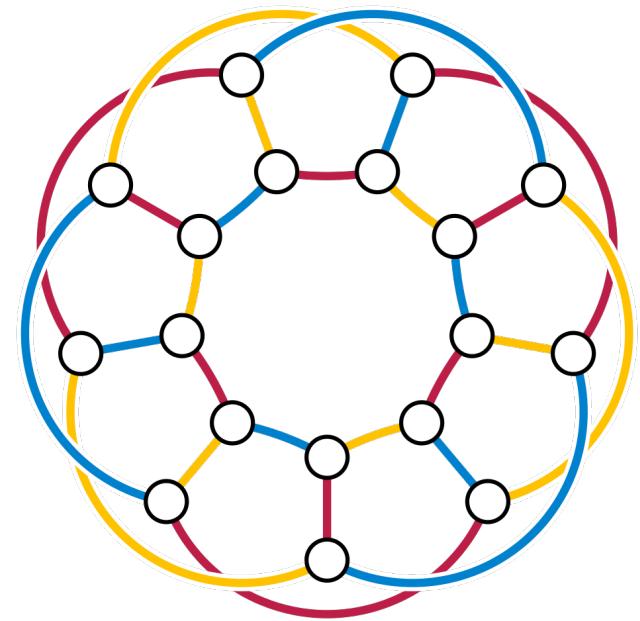
- This restriction is stricter on **dense** graphs and less strict on **sparse** graphs.
- The standard variants are either too trivial or too hard on sparse graphs.
- Similarly, the **Semi**-MPC model is also defined.

# Edge Coloring

Streaming and Massively Parallel Algorithms for Edge Coloring [**ESA'19**]  
Behnezhad, Derakhshan, Knittel, Hajiaghayi, **me**

# The Edge Coloring problem

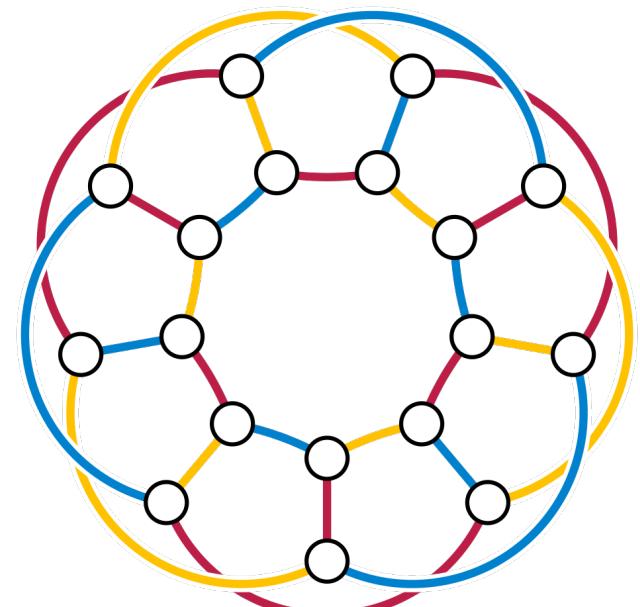
- Given a graph  $G = (V, E)$ , a **valid** “edge-coloring” is a function  $\text{COL}: E \rightarrow [\Psi]$  so that no two **incident edges** have a common color.
- We wish to minimize  $\Psi$ , i.e., the number of colors.



# The Edge Coloring problem

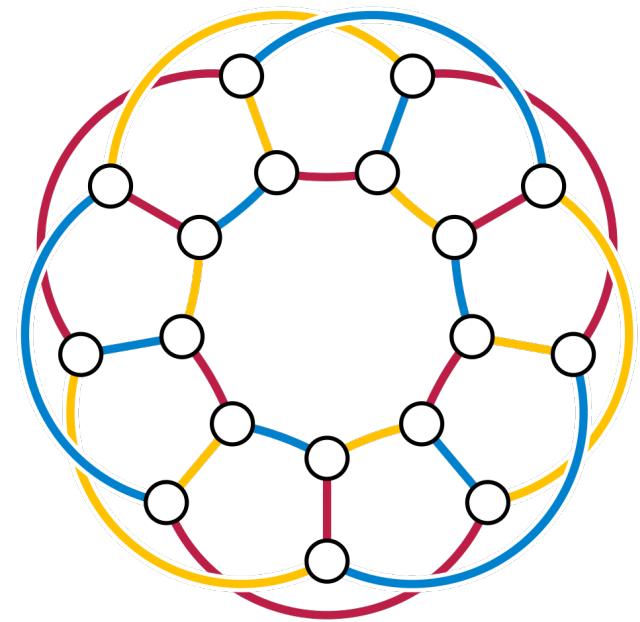
- Given a graph  $G = (V, E)$ , a **valid** “edge-coloring” is a function  $\text{COL}: E \rightarrow [\Psi]$  so that no two **incident edges** have a common color.
- Let  $\Delta$  be the maximum degree of graph  $G$ .
- Then, we know  $\Delta \leq \Psi \leq \Delta + 1$ .
- An odd cycle ( $\Delta = 2$ ) needs **3** colors.

[Vizing 1964]



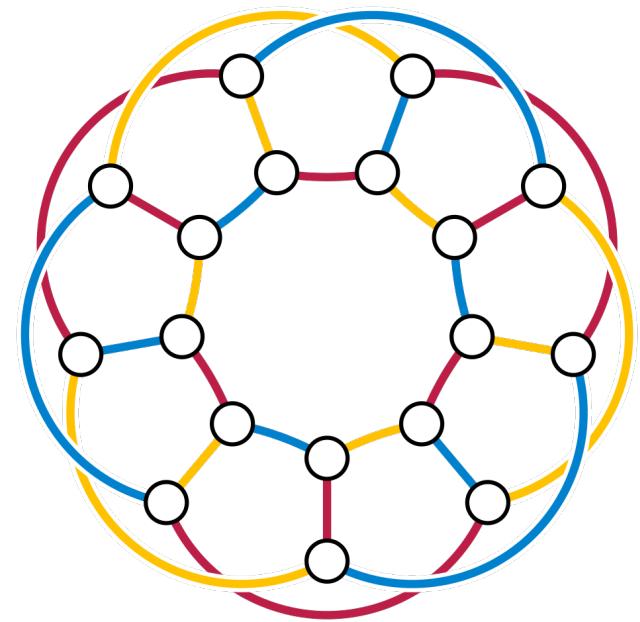
# The Edge Coloring problem

- Given a graph  $G = (V, E)$ , a **valid** “edge-coloring” is a function  $\text{COL}: E \rightarrow [\Psi]$  so that no two **incident edges** have a common color.
- However,  $(\Delta + 1)$ -coloring algorithms are highly **sequential**.
- There is a greedy  $(2\Delta - 1)$ -coloring algorithm.



# The Edge Coloring problem

- Given a graph  $G = (V, E)$ , a **valid** “edge-coloring” is a function  $\text{COL}: E \rightarrow [\Psi]$  so that no two **incident edges** have a common color.
- There is a greedy  $(2\Delta - 1)$ -coloring algorithm.
  - Process edges in an arbitrary order.
  - Color each edge with an available color.



# Related work

- The closely related **Vertex** Coloring problem has constant-round MPC and one-pass Streaming algorithms both with  $S = \tilde{O}(|V|)$  and  $(\Delta + 1)$ -colors used.[Assadi et. al 2019]
- Unfortunately **not applicable** to Edge Coloring.
- Differences:
  - Similar **greedy** algorithm gives a  $(\Delta + 1)$ -vertex-coloring.
  - We can store the **output** inside the machines.

# Related work

- The only previously known **MPC** algorithm for **Edge** Coloring needs **superlinear** space per machine to achieve a  $(\Delta + o(\Delta))$ -coloring. [Harvey et. al 2018]
- Little to no result known for **Edge** Coloring in the **Streaming** model.

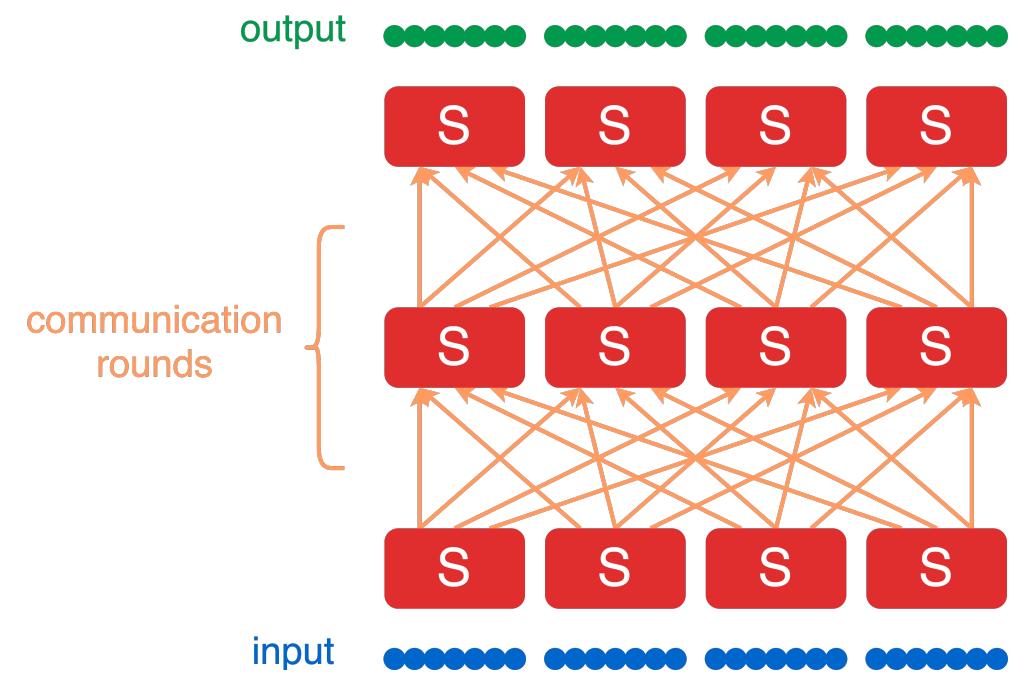
# Edge Coloring

Massively Parallel Computation

# MPC Edge Coloring

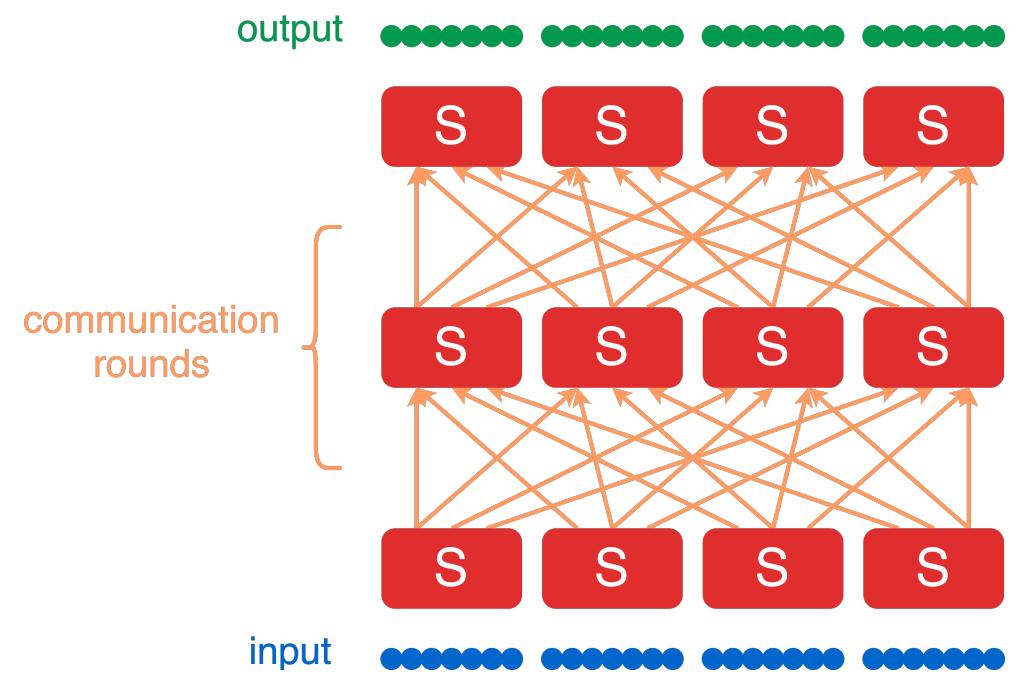
- A graph  $G = (V, E)$  is given where  $|V| = n$  and  $|E| = m$ , i.e.  $N = O(n + m)$ .
- There is a **constant**-round MPC algorithm, with  $S = O(n)$  and  $S \cdot M = O(m)$ , which computes an edge-coloring so that  $\Psi = \Delta + \tilde{O}(\Delta^{3/4})$ .

[BDKHS 2019]



# Semi-MPC Edge Coloring

- A graph  $G = (V, E)$  is given where  $|V| = n$  and  $|E| = m$ , i.e.  $N = O(n + m)$ .
- There is a **constant**-round MPC algorithm, with  $S = O(n)$  and  $S \cdot M = O(m)$ , which computes an edge-coloring so that  $\Psi = \Delta + \tilde{O}(\Delta^{3/4})$ .  
[BDKHS 2019]
- It is technically a **Semi**-MPC algorithm.

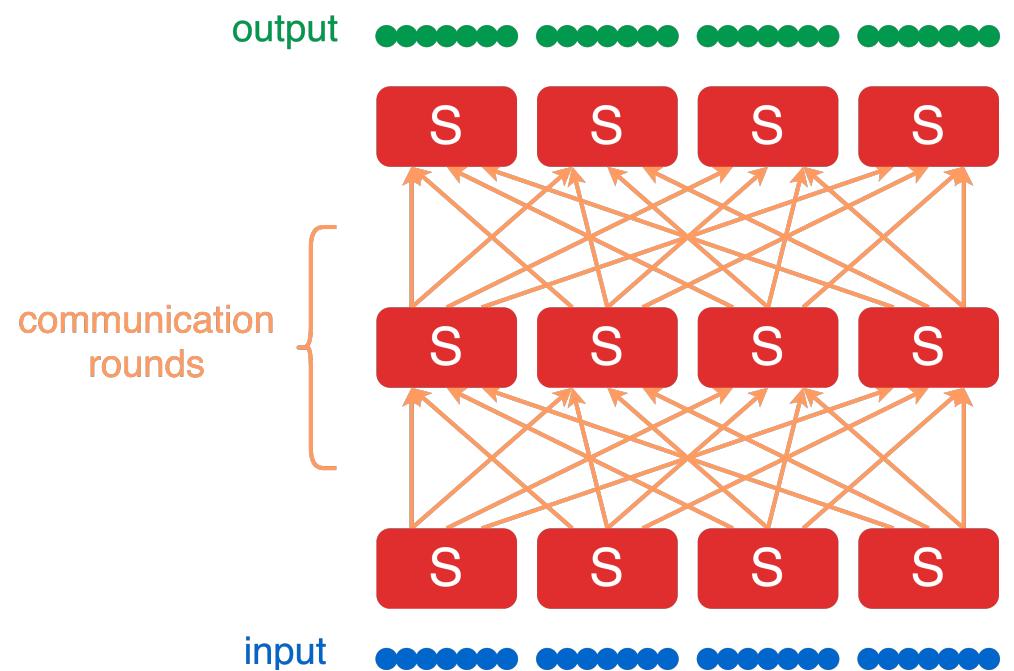


# Semi-MPC Edge Coloring

- It is technically a **Semi**-MPC algorithm, but we can achieve  $o(n)$  space in dense graphs.
- The exact space-per-machine is equal to

$$S = O\left(\frac{n\Delta}{k^2} + \frac{n}{k}\sqrt{\frac{\Delta}{k}\log n}\right)$$

- Set  $k = \sqrt{\Delta} + \log n$ .

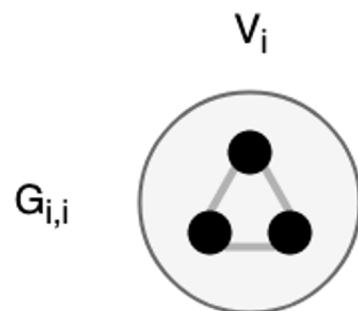


# Vertex Partitioning

- Set  $\textcolor{brown}{k} = \sqrt{\Delta} + \log n$ .
- Random vertex partitioning:  $V = V_1 \cup V_2 \cup \dots \cup V_{\textcolor{brown}{k}}$ .

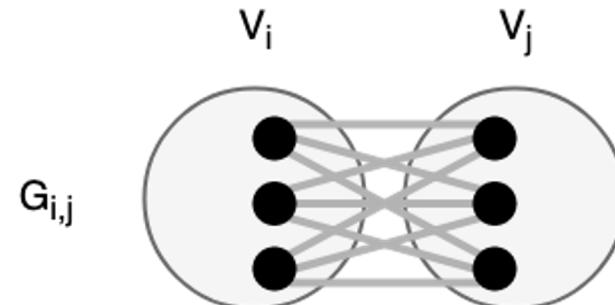
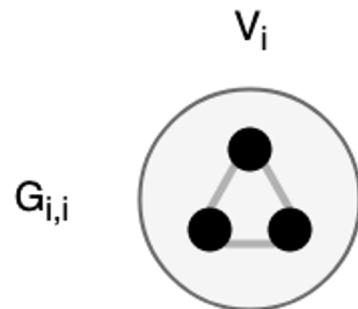
# Vertex Partitioning

- Set  $\textcolor{brown}{k} = \sqrt{\Delta} + \log n$ .
- Random vertex partitioning:  $V = V_1 \cup V_2 \cup \dots \cup V_{\textcolor{brown}{k}}$ .
- $G_{i,i} = \{ (u,v) \mid u \in V_i \wedge v \in V_i \}$ : the **induced** subgraph of each partition



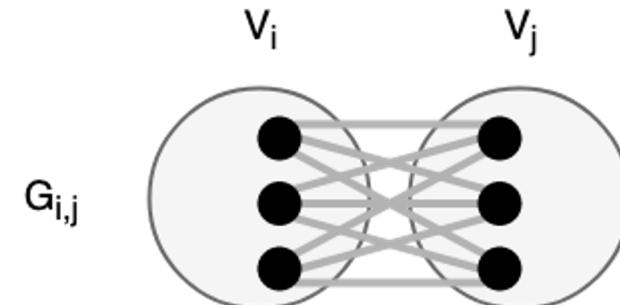
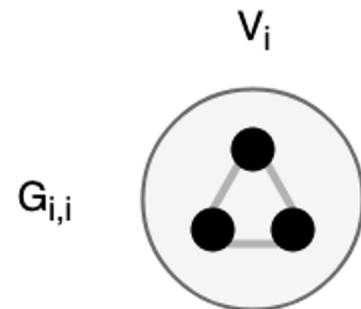
# Vertex Partitioning

- Random vertex partitioning:  $V = V_1 \cup V_2 \cup \dots \cup V_k$ .
- $G_{i,i} = \{ (u, v) \mid u \in V_i \wedge v \in V_i \}$ : the **induced** subgraph of each partition
- $G_{i,j} = \{ (u, v) \mid u \in V_i \wedge v \in V_j \}$ : the **bipartite induced** subgraph of pairs of partitions



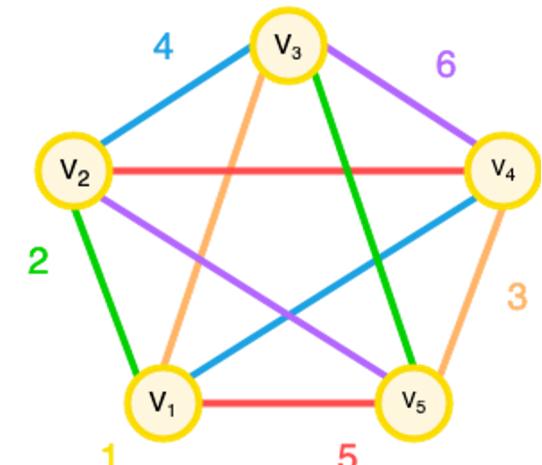
# Local Coloring

- $G_{i,i} = \{ (u, v) \mid u \in V_i \wedge v \in V_i \}$ : the **induced** subgraph of each partition
- $G_{i,j} = \{ (u, v) \mid u \in V_i \wedge v \in V_j \}$ : the **bipartite induced** subgraph of pairs of partitions
- Run Vizing's  $(\Delta + 1)$ -coloring algorithm on each machine.



# Merging Colored Subgraphs

- $k + 1$  disjoint color **palettes** are enough.
- The number of colors in each **palette** needs to be as large as the maximum degree in subgraphs.
- The maximum degree in each subgraphs is **concentrated**.
- For  $k = \sqrt{\Delta} + \log n$ , we have  $\Psi = \Delta + \tilde{O}(\Delta^{3/4})$ .



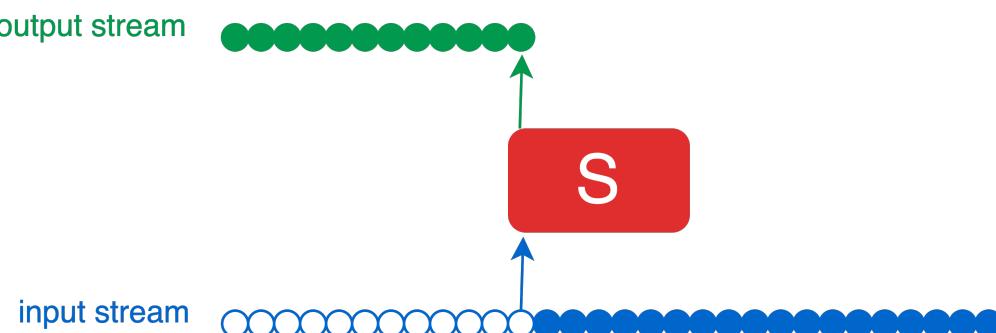
# **Edge Coloring**

## **Random Streaming**

# Random Streaming Edge Coloring

- A graph  $G = (V, E)$  is given where  $|V| = n$  and  $|E| = m$ , i.e.  $\textcolor{brown}{N} = O(n + m)$ .
- There is a **one**-pass  $\textcolor{brown}{W}$ -streaming algorithm, with  $\textcolor{brown}{S} = O(n)$ , which streams a valid edge-coloring so that  $\textcolor{brown}{\Psi} = (2e + \epsilon)\Delta$ .

[BDKHS 2019]

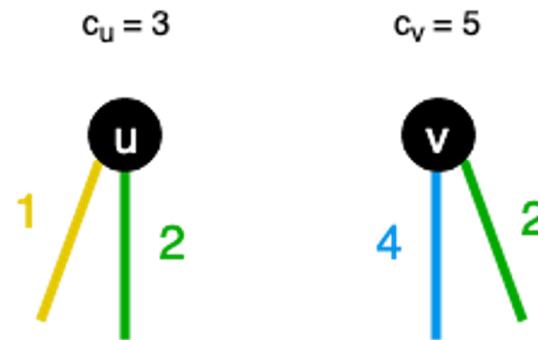


# Random Streaming Edge Coloring

- A graph  $G = (V, E)$  is given where  $|V| = n$  and  $|E| = m$ , i.e.  $N = O(n + m)$ .
- There is a **one**-pass **W**-streaming algorithm, with  $S = O(n)$ , which streams a valid edge-coloring so that  $\Psi = (2e + \epsilon)\Delta$ .[BDKHS 2019]
- The algorithm is straight-forward.
- Maintain a **counter** variable  $c_v$  for each vertex, initially set to 1.

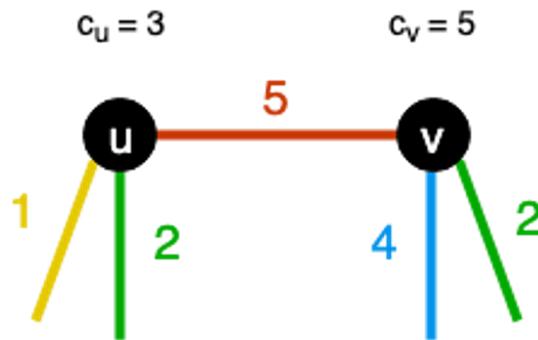
# Random Streaming Edge Coloring

- Maintain a **counter** variable  $c_v$  for each vertex, initially set to 1.
- Upon arrival of  $(u, v)$  color it  $\max(c_v, c_u)$ .



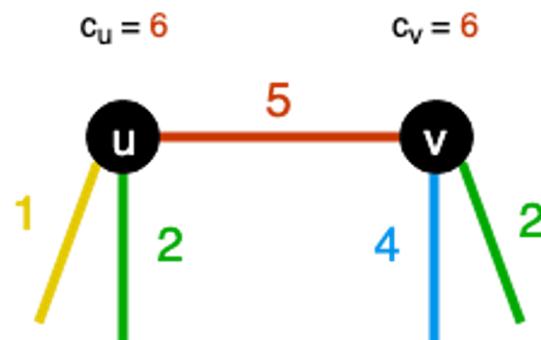
# Random Streaming Edge Coloring

- Maintain a **counter** variable  $c_v$  for each vertex, initially set to 1.
- Upon arrival of  $(u, v)$  color it  $\max(c_v, c_u)$ .



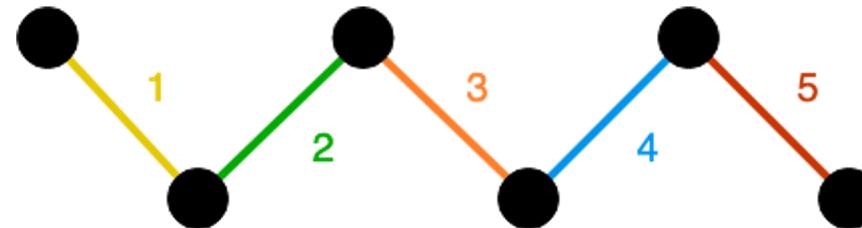
# Random Streaming Edge Coloring

- Maintain a **counter** variable  $c_v$  for each vertex, initially set to 1.
- Upon arrival of  $(u, v)$  color it  $\max(c_v, c_u)$ .
- Set both  $c_v$  and  $c_u$  to  $\max(c_v, c_u) + 1$ .



# Longest **Monotone** Path

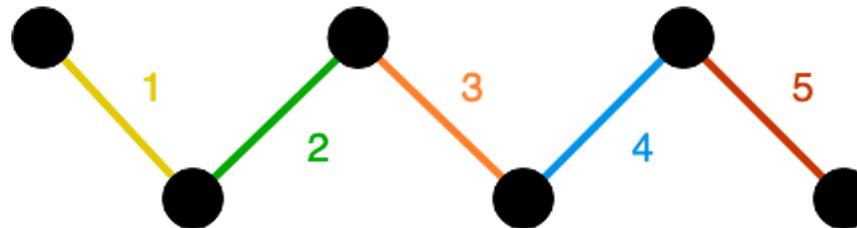
- o **Lemma:**  $\Psi$  is equal to the length of the longest **monotone** path in the **line graph** after the algorithm finishes.
- o We can construct a **monotone** path of length  $\Psi$  starting from an edge colored  $\Psi$ .



# Longest **Monotone** Path

- o **Lemma:**  $\Psi$  is equal to the length of the longest **monotone** path in the **line graph** after the algorithm finishes.

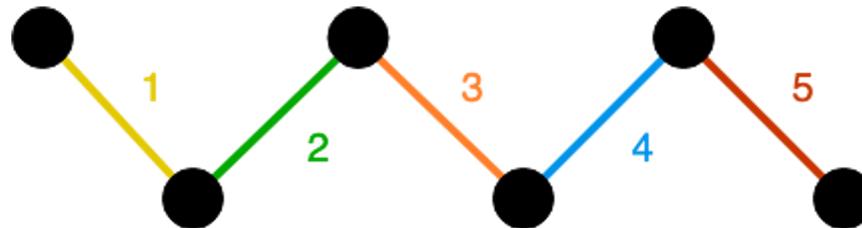
$$\Pr[\Psi \geq \alpha\Delta] \leq \frac{(2\Delta)^{\alpha\Delta}}{(\alpha\Delta)!}$$



# Longest **Monotone** Path

- o **Lemma:**  $\Psi$  is equal to the length of the longest **monotone** path in the **line graph** after the algorithm finishes.

$$\Pr[\Psi \geq (2e + \epsilon)\Delta] \leq n^{-c}$$



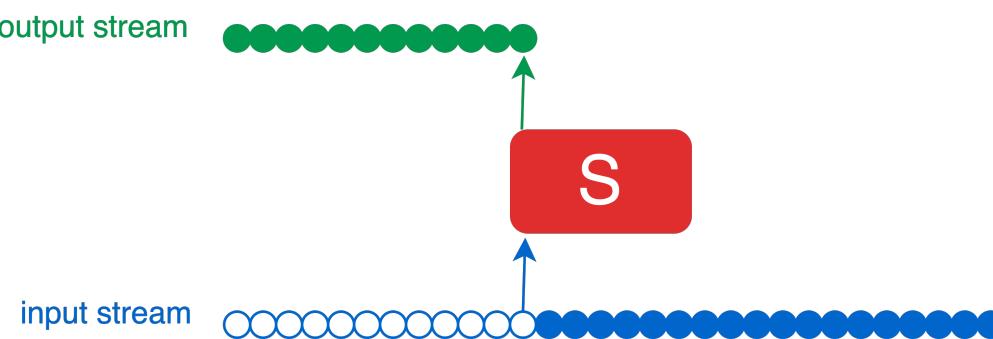
# **Edge Coloring**

## **Adversarial** Streaming

# Adversarial Streaming Edge Coloring

- A graph  $G = (V, E)$  is given where  $|V| = n$  and  $|E| = m$ , i.e.  $\textcolor{brown}{N} = O(n + m)$ .
- There is a **one**-pass  $\textcolor{brown}{W}$ -streaming algorithm, with  $\textcolor{brown}{S} = O(n)$ , which streams a valid edge-coloring so that  $\textcolor{brown}{\Psi} = O(\Delta^2)$ .

[BDKHS 2019]

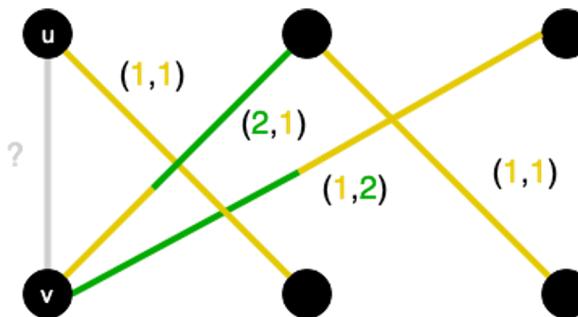


# Adversarial Streaming Edge Coloring

- A graph  $G = (V, E)$  is given where  $|V| = n$  and  $|E| = m$ , i.e.  $\textcolor{brown}{N} = O(n + m)$ .
- There is a **one**-pass  $\textcolor{brown}{W}$ -streaming algorithm, with  $\textcolor{brown}{S} = O(n)$ , which streams a valid edge-coloring so that  $\textcolor{brown}{\Psi} = O(\Delta^2)$ .**[BDKHS 2019]**
- The algorithm is very similar to the random stream algorithm.
- Maintain a **counter** variable  $c_v$  for each vertex, initially set to 1.

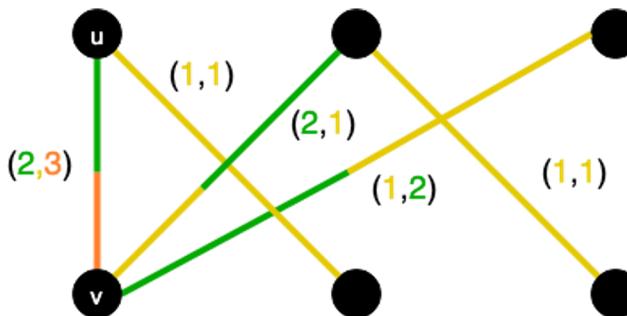
# Bipartite Edge Coloring

- Maintain a **counter** variable  $c_v$  for each vertex, initially set to 1. Also assume the graph is **bipartite**.
- Upon arrival of  $(u, v)$  color it  $(c_u, c_v)$ .



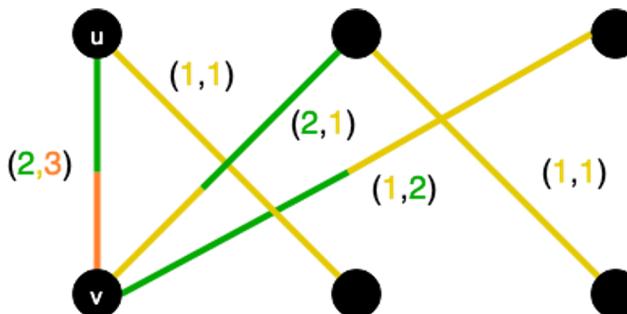
# Bipartite Edge Coloring

- Maintain a **counter** variable  $c_v$  for each vertex, initially set to 1. Also assume the graph is **bipartite**.
- Upon arrival of  $(u, v)$  color it  $(c_u, c_v)$ .



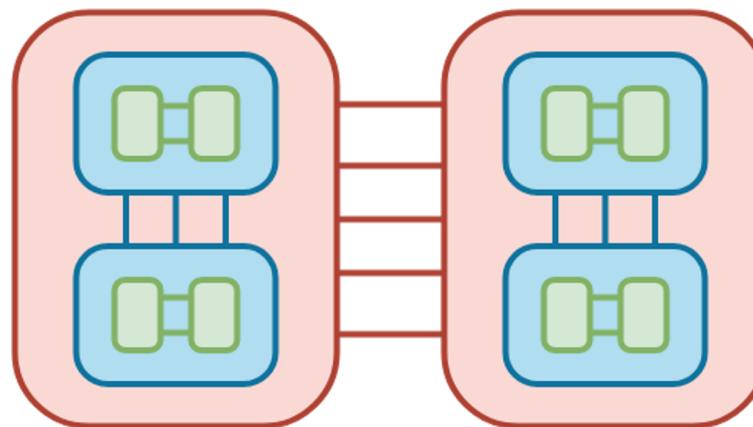
# Bipartite Edge Coloring

- Maintain a **counter** variable  $c_v$  for each vertex, initially set to **1**. Also assume the graph is **bipartite**.
- Upon arrival of  $(u, v)$  color it  $(c_u, c_v)$ .
- Increase both  $c_u$  and  $c_v$  by **1**.



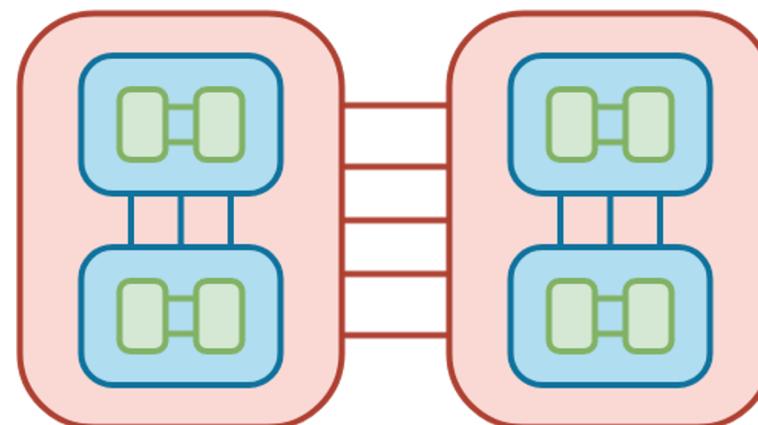
# General Edge Coloring

- A **bipartite** graph is colored with  $\Delta^2$  colors using this method.
- Decompose the graph into  $O(\log n)$  random bipartite graphs, and color each with a different **palette**.



# General Edge Coloring

- A **bipartite** graph is colored with  $\Delta^2$  colors using this method.
- Decompose the graph into  $O(\log n)$  random bipartite graphs, and color each with a different **palette**.
- In total,  $\Psi = O(\Delta^2)$ .



**Thanks for watching!**

Any questions?